

[Click here to view the PDF version.](#)

```
1 begin
2   using LinearAlgebra
3   using SparseArrays
4   using PlutoUI
5   using PlutoTeachingTools
6   using HypertextLiteral: @html, @html_str
7 end
```

≡ Table of Contents

Direct methods for linear systems

A difficult example

Motivation: The \ (backslash) operator

Solving triangular systems

Gaussian elimination and LU factorisation

Running Algorithm 3

LU factorisation with pivoting

Solving linear systems based on LU factorisation

Conclusion: What backslash \ does to solve

Optional: More details on pivoting

Memory usage and fill-in

Sparse matrices

Memory: LU factorisation of full matrices

Memory: LU factorisation of sparse matrices

Computational cost of LU factorisation

Scalar product

Matrix-vector product

LU factorisation

Numerical stability

Matrix condition numbers and stability result

Computing condition numbers

Direct methods for linear systems

In the previous chapter on polynomial interpolation we were already confronted with the need to solve linear systems, that is a system of equations of the form

$$\mathbf{Ax} = \mathbf{b}, \tag{1}$$

where we are given a matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ as well as a right-hand side $\mathbf{b} \in \mathbb{R}^n$. As the solution we seek the unknown $\mathbf{x} \in \mathbb{R}^n$.

A difficult example

Solving linear equations is a standard exercise in linear algebra and you probably have already done it in previous courses. However, you might also know that solving such linear systems is not always equally easy.

This one, for example, is particularly difficult:

```
M_difficult = 2×2 Matrix{Float64}:  
  3.14159  2.71828  
  3.14159  2.71828
```

```
1 M_difficult = [π      e;  
2               355/113  23225/8544] # Rational approximations to π and e
```

```
b_difficult = ▶ [1.0, 1.0]
```

```
1 b_difficult = [1.0, 1.0]
```

Solving it in 64-bit Float64 gives this result:

```
▶ [-0.00958449, 0.378957]
```

```
1 M_difficult \ b_difficult
```

while using 32-bit Float32 gives this:

```
▶ [-0.114318, 0.5]
```

```
1 let  
2   M_32 = convert(Matrix{Float32}, M_difficult)  
3   b_32 = convert(Vector{Float32}, b_difficult)  
4   M_32 \ b_32  
5 end
```

These clearly differ and we would probably trust the `Float64` result more. But is this justified? And what are conditions to quantify that a linear system is more challenging than another?

This is what we want to explore in this part of the course.

Motivation: The `\` (backslash) operator

For solving linear systems, we so far employed Julia's backslash `\` operator. So let's find out what it actually does under the hood. We take the problem

```
1 begin
2     A = Float64[-4  3 -1;
3                 2  1  0;
4                 4 -3  4]
5     b = [2, 4, -2]
6 end;
```

as an example. Normally we would now just do `A \ b`, so let's check what this calls

1 method for generic function `\` from Base:

- `\(A::AbstractMatrix, B::AbstractVecOrMat)` in LinearAlgebra at [/opt/hostedtoolcache/julia/1.11.4/x64/share/julia/stdlib/v1.11/LinearAlgebra/src/generic.jl:1118](https://github.com/JuliaLang/julia/blob/v1.11.4/src/LinearAlgebra/src/generic.jl#L1118)

```
1 methods(\, (Matrix, Vector))
```

Ok, so this calls into Julia's linear algebra library. The code is [located here](#). Essentially it performs an **LU factorisation**:

```
LU = LinearAlgebra.LU{Float64, Matrix{Float64}, Vector{Int64}}
L factor:
3×3 Matrix{Float64}:
 1.0  0.0  0.0
-0.5  1.0  0.0
-1.0  0.0  1.0
U factor:
3×3 Matrix{Float64}:
-4.0  3.0 -1.0
 0.0  2.5 -0.5
 0.0  0.0  3.0
```

```
1 LU = lu(A)
```

which as we can see *factorises* the matrix **A** into a **lower triangular** matrix **L** and an **upper triangular** matrix **U**. More formally:

Definition: LU factorisation

Given a matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ find a lower triangular matrix $\mathbf{L} \in \mathbb{R}^{n \times n}$ and an upper triangular matrix $\mathbf{U} \in \mathbb{R}^{n \times n}$, such that

$$\mathbf{A} = \mathbf{L}\mathbf{U} \quad (2)$$

We can easily check that by multiplying $\mathbf{L}$ and $\mathbf{U}$ we indeed recover $\mathbf{A}$:

```
3x3 Matrix{Float64}:  
-4.0  3.0 -1.0  
 2.0  1.0  0.0  
 4.0 -3.0  4.0
```

```
1 LU.L * LU.U
```

```
3x3 Matrix{Float64}:  
0.0  0.0  0.0  
0.0  0.0  0.0  
0.0  0.0  0.0
```

```
1 A - LU.L * LU.U
```

Now we are a step closer to what happens, but why is this useful ?

Solving triangular systems

The factorisation into two triangular matrices is useful, because it is especially easy to solve a system where the matrix is triangular. For example consider the *lower* triangular system

$$\begin{pmatrix} 4 & 0 & 0 & 0 \\ 3 & -1 & 0 & 0 \\ -1 & 0 & 3 & 0 \\ 1 & -1 & -1 & 2 \end{pmatrix} \mathbf{x} = \begin{pmatrix} 8 \\ 5 \\ 0 \\ 1 \end{pmatrix}.$$

- The first row of this system simply states $4x_1 = 8$, which is very easy to solve, namely $x_1 = 8/4 = 2$.
- The second row states $3x_1 - x_2 = 5$. However, x_1 is already known and can be inserted to find $x_2 = -(5 - 3 \cdot 2) = 1$.
- Following the same idea the third row gives $x_3 = (0 + 1 \cdot 2)/3 = 2/3$ and the last row $x_4 = (1 - 1 \cdot 2 + 1 \cdot 1 + 1 \cdot 2/3)/2 = 1/3$.
- In total we found

$$\mathbf{x} = \begin{pmatrix} 2 \\ 1 \\ 2/3 \\ 1/3 \end{pmatrix}.$$

Generalising to an arbitrary 4x4 lower-rectangular matrix

$$\begin{pmatrix} L_{11} & 0 & 0 & 0 \\ L_{21} & L_{22} & 0 & 0 \\ L_{31} & L_{32} & L_{33} & 0 \\ L_{41} & L_{42} & L_{43} & L_{44} \end{pmatrix} \mathbf{x} = \begin{pmatrix} b_1 \\ b_2 \\ b_3 \\ b_4 \end{pmatrix}$$

this solution algorithm can be expressed as

$$\begin{aligned} x_1 &= \frac{b_1}{L_{11}}, \\ x_2 &= \frac{b_2 - L_{21}x_1}{L_{22}}, \\ x_3 &= \frac{b_3 - L_{31}x_1 - L_{32}x_2}{L_{33}}, \\ x_4 &= \frac{b_4 - L_{41}x_1 - L_{42}x_2 - L_{43}x_3}{L_{44}}. \end{aligned} \tag{3}$$

We obtain

Algorithm 1: Forward substitution

Given a lower-triangular matrix $\mathbf{L} \in \mathbb{R}^{n \times n}$ a linear system $\mathbf{L}\mathbf{x} = \mathbf{b}$ can be solved by looping from $i = 1, \dots, n$ and computing

$$x_i = \frac{1}{L_{ii}} \left(b_i - \sum_{j=1}^{i-1} L_{ij}x_j \right)$$

or in form of an implementation

forward_substitution (generic function with 1 method)

```
1 function forward_substitution(L, b)
2     n = size(L, 1)
3     x = zeros(n)
4     x[1] = b[1] / L[1, 1]
5     for i in 2:n
6         row_sum = 0.0
7         for j in 1:i-1
8             row_sum += L[i, j] * x[j]
9         end
10
11         x[i] = 1 / L[i, i] * (b[i] - row_sum)
12     end
13     x
14 end
```

For upper triangular matrices we solve linear systems proceeds following the same idea — just in this case we start from the last row and not the first. For example a system

$$\begin{pmatrix} U_{11} & U_{12} & U_{13} & U_{14} \\ 0 & U_{22} & U_{23} & U_{24} \\ 0 & 0 & U_{33} & U_{34} \\ 0 & 0 & 0 & U_{44} \end{pmatrix} \mathbf{x} = \begin{pmatrix} b_1 \\ b_2 \\ b_3 \\ b_4 \end{pmatrix}$$

we solve by starting with x_4 and working our way forward to x_1 :

$$\begin{aligned} x_4 &= \frac{b_4}{U_{44}}, \\ x_3 &= \frac{b_3 - U_{34}x_4}{U_{33}}, \\ x_2 &= \frac{b_2 - U_{23}x_3 - U_{24}x_4}{U_{22}}, \\ x_1 &= \frac{b_1 - U_{12}x_2 - U_{13}x_3 - U_{14}x_4}{U_{11}}. \end{aligned} \tag{4}$$

More formally this algorithm is called

Algorithm 2: Backward substitution

Given an upper-triangular matrix $\mathbf{U} \in \mathbb{R}^{n \times n}$ a linear system $\mathbf{U}\mathbf{x} = \mathbf{b}$ can be solved by looping in reverse order $i = n, \dots, 1$ and computing

$$x_i = \frac{1}{U_{ii}} \left(b_i - \sum_{j=i+1}^n U_{ij} x_j \right)$$

or in code

backward_substitution (generic function with 1 method)

```

1 function backward_substitution(U, b)
2     n = size(U, 1)
3     x = zeros(n)
4     x[n] = b[n] / U[n, n]
5     for i in n-1:-1:1 # Note that this loop goes *backwards* n-1, n-2, ..., 1
6         row_sum = 0.0
7         for j in i+1:n
8             row_sum += U[i, j] * x[j]
9         end
10        x[i] = (b[i] - row_sum) / U[i, i]
11    end
12    x
13 end

```

Based on our discussion we now understand why it is advantageous to perform an LU factorisation when solving a linear system: For both the resulting triangular matrix $\mathbf{L}$ as well as the matrix $\mathbf{U}$, simple solution algorithms are available. In summary we thus obtain

Algorithm 3: Solving linear systems by LU factorisation / Gaussian elimination

We are given a matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$, a right-hand side $\mathbf{b} \in \mathbb{R}^n$. This algorithm computes the solution $\mathbf{x} \in \mathbb{R}^n$ to $\mathbf{Ax} = \mathbf{b}$:

1. Factorise $\mathbf{A} = \mathbf{LU}$.
2. Solve $\mathbf{Lz} = \mathbf{b}$ for $\mathbf{z}$ using **forward substitution** (Algorithm 1).
3. Solve $\mathbf{Ux} = \mathbf{z}$ for $\mathbf{x}$ using **backward substitution** (Algorithm 2).

A few remarks are in order:

- We have so far not discussed how to even compute the LU factorisation. As we will see in the next section this step will actually be accomplished by the **Gaussian elimination algorithm**, which probably already know from your linear algebra lectures.
- A severe **drawback of this naive algorithm** is immediately apparent from our 4x4 example problem, where (3) and (4) provide explicit solution expressions. If by any chance an

element L_{ii} or U_{ii} happens to be zero, the algorithm cannot work. So one needs to come up with ways around this. We will mention a few ideas in the section on **pivoting**.

- Note, that **steps 2. and 3.** of the algorithm **do not depend on the original matrix $\mathbf{A}$** , but only on the right hand side $\mathbf{b}$. In other words once the factorisation $\mathbf{A} = \mathbf{LU}$ has been computed, solving $\mathbf{Ax} = \mathbf{b}$ for different right hand sides $\mathbf{b}$ only requires the execution of steps 2. and 3. As we will see later, step 1. is the most expensive step. Therefore **once $\mathbf{A}$ has been factorised**, the **cost for solving an arbitrary linear system** involving $\mathbf{A}$ is reduced as its factorised form $\mathbf{LU}$ can be used in its place.

Gaussian elimination and LU factorisation

We will now discuss how to compute an LU factorization. As it turns out the Gaussian elimination algorithm you already learned in previous linear algebra classes provides exactly this missing ingredient.

Recall that **Gaussian elimination** allows us to transform a linear system into an upper triangular system using elementary operations, which combine the different equations.

Example: Manual Gaussian elimination

We will discuss one example, namely the linear system

$$\underbrace{\begin{pmatrix} 2 & 1 & 0 \\ -4 & 3 & -1 \\ 4 & -3 & 4 \end{pmatrix}}_{=\mathbf{A}} \underbrace{\begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix}}_{=\mathbf{x}} = \underbrace{\begin{pmatrix} 4 \\ 2 \\ -2 \end{pmatrix}}_{=\mathbf{b}}.$$

Let us call $r_1^{(1)}$, $r_2^{(1)}$ and $r_3^{(1)}$ the three equations of the system:

$$\begin{array}{lllll} r_1^{(1)} : & 2x_1 & +x_2 & & =4 \\ r_2^{(1)} : & -4x_1 & +3x_2 & -x_3 & =2 \\ r_3^{(1)} : & 4x_1 & -3x_2 & +4x_3 & =-2 \end{array}$$

on which we perform the following operations.

1st step: Zero out all but the first row in the first column.

$$\begin{aligned}
r_1^{(2)} &\leftarrow r_1^{(1)} && \Rightarrow 2x_1 & +x_2 & & =4 \\
r_2^{(2)} &\leftarrow r_2^{(1)} - \underbrace{\left(\frac{-4}{2}\right)r_1^{(1)}}_{L_{21}} && \Rightarrow & 5x_2 & -x_3 & =10 \\
r_3^{(2)} &\leftarrow r_3^{(1)} - \underbrace{\left(\frac{4}{2}\right)r_1^{(1)}}_{L_{31}} && \Rightarrow & -5x_2 & +4x_3 & =-10.
\end{aligned}$$

2nd step: Zero out all but the first and second row in the second column.

$$\begin{aligned}
r_1^{(3)} &\leftarrow r_1^{(2)} && \Rightarrow 2x_1 & +x_2 & & =4 \\
r_2^{(3)} &\leftarrow r_2^{(2)} && \Rightarrow & 5x_2 & -x_3 & =10 \\
r_3^{(3)} &\leftarrow r_3^{(2)} - \underbrace{\left(\frac{-5}{5}\right)r_2^{(2)}}_{L_{32}} && \Rightarrow & & 3x_3 & =0.
\end{aligned}$$

After this step we have thus obtained an upper-triangular system, which we would solve using backward substitution. Note that during the iterations the system matrix got slowly reduced to an upper triangular matrix, which we can identify with **U**:

$$\underbrace{\mathbf{A}^{(1)}}_{=\mathbf{A}} = \begin{pmatrix} 2 & 1 & 0 \\ -4 & 3 & -1 \\ 4 & -3 & 4 \end{pmatrix} \Rightarrow \mathbf{A}^{(2)} = \begin{pmatrix} 2 & 1 & 0 \\ 0 & 5 & -1 \\ 0 & -5 & 4 \end{pmatrix} \Rightarrow \mathbf{A}^{(3)} = \begin{pmatrix} 2 & 1 & 0 \\ 0 & 5 & -1 \\ 0 & 0 & 3 \end{pmatrix} = \mathbf{U}$$

In turn the matrix **L** is obtained by collecting the factors we employed to multiply the rows. We add diagonal entries of **1** to indicate that the *i*-th row has been left unchanged in the *i*-th step. Accumulating step by step we obtain

$$\mathbf{L}^{(1)} = \begin{pmatrix} 1 & & \\ -\frac{4}{2} & 1 & \\ \frac{4}{2} & & 1 \end{pmatrix} \Rightarrow \mathbf{L}^{(2)} = \begin{pmatrix} 1 & & \\ -2 & 1 & \\ 2 & -\frac{5}{5} & 1 \end{pmatrix} \Rightarrow \mathbf{L}^{(3)} = \begin{pmatrix} 1 & & \\ -2 & 1 & \\ 2 & -1 & 1 \end{pmatrix} = \mathbf{L}$$

Finally by multiplying out the matrices we easily verify

$$\mathbf{LU} = \begin{pmatrix} 1 & 0 & 0 \\ -2 & 1 & 0 \\ 2 & -1 & 1 \end{pmatrix} \begin{pmatrix} 2 & 1 & 0 \\ 0 & 5 & -1 \\ 0 & 0 & 3 \end{pmatrix} = \begin{pmatrix} 2 & 1 & 0 \\ -4 & 3 & -1 \\ 4 & -3 & 4 \end{pmatrix} = \mathbf{A}$$

So effectively Gaussian elimination provides us with an approach to compute the LU factorisation. As a result Algorithm 3 is nothing but a formalised form of the procedure you used

in your linear algebra lectures to solve linear systems using Gaussian elimination. We summarise:

Algorithm 4: LU factorisation

Input: $\mathbf{A} \in \mathbb{R}^{n \times n}$

Output: $\mathbf{U} \in \mathbb{R}^{n \times n}$, $\mathbf{L} \in \mathbb{R}^{n \times n}$

- $\mathbf{A}^{(1)} = \mathbf{A}$
- for $k = 1, \dots, n - 1$ (algorithm steps)
 - $L_{kk} = 1$
 - for $i = k + 1, \dots, n$ (Loop over rows)
 - $L_{ik} = \frac{A_{ik}^{(k)}}{A_{kk}^{(k)}}$
 - for $j = k, \dots, n$ (Loop over columns)
 - $A_{ij}^{(k+1)} = A_{ij}^{(k)} - L_{ik}A_{kj}^{(k)}$
- $L_{nn} = 1$ (the loop above only runs until $n - 1$)
- $\mathbf{U} = \mathbf{A}^{(n)}$

An implementation of LU factorisation could be realised as such:

factorise_lu (generic function with 1 method)

```
1 function factorise_lu(A)
2     n = size(A, 1)
3     L = zeros(n, n)    # Initialise L and U by zeros
4     Ak = float(copy(A)) # Make a copy of A and ensure that all entries
5                          # are converted to floating-point numbers
6
7     for k in 1:n-1      # Algorithm steps
8         L[k, k] = 1.0
9         for i in k+1:n   # Loop over rows
10            L[i, k] = Ak[i, k] / Ak[k, k]
11            for j in k:n  # Loop over columns
12                # Update Ak in-place to give Ak+1
13                Ak[i, j] = Ak[i, j] - L[i, k] * Ak[k, j]
14            end
15        end
16    end
17    L[n, n] = 1.0        # Since the loop only runs until n-1
18    U = Ak
19
20    # Return L and U using Julia datastructures to indicate
21    # their special lower-triangular and upper-triangular form.
22    return LowerTriangular(L), UpperTriangular(U)
23 end
```

We stay with the example where we performed manual Gaussian elimination:

```
A_manual = 3×3 Matrix{Float64}:
 2.0  1.0  0.0
-4.0  3.0 -1.0
 4.0 -3.0  4.0
```

```
1 A_manual = Float64[ 2  1  0;
2                     -4  3 -1;
3                     4 -3  4]
```

With this slider you can stop factorise_lu after 1, 2 or 3 steps, checking that agrees with the steps we computed manually:

• nstep_lu_A =  0

```
► (Ak = 3×3 Matrix{Float64}::, L = 3×3 Matrix{Float64}::)
 2.0  1.0  0.0      0.0  0.0  0.0
```

Running Algorithm 3

With this we have the missing ingredient to run Algorithm 3 and numerically solve a linear system.

We stay with the problem

$$\underbrace{\begin{pmatrix} 2 & 1 & 0 \\ -4 & 3 & -1 \\ 4 & -3 & 4 \end{pmatrix}}_{=\mathbf{A}_{\text{manuel}}} \underbrace{\begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix}}_{=\mathbf{x}} = \underbrace{\begin{pmatrix} 4 \\ 2 \\ -2 \end{pmatrix}}_{=\mathbf{b}_{\text{manuel}}}$$

which we solved manually beforehand:

```
3x3 Matrix{Float64}:  
 2.0  1.0  0.0  
-4.0  3.0 -1.0  
 4.0 -3.0  4.0
```

```
1 A_manual
```

```
b_manual = ▶ [4, 2, -2]
```

```
1 b_manual = [4, 2, -2]
```

We follow **Algorithm 3**. First we need to find the LU factorisation:

```
▶ (3x3 LowerTriangular{Float64, Matrix{Float64}}:, 3x3 UpperTriangular{Float64, Matrix{  
 1.0 . . 2.0 1.0 0.0
```

```
1 L, U = factorise_lu(A_manual)
```

which agrees what we obtained manually. Next we forward substitute:

```
z = ▶ [4.0, 10.0, 0.0]
```

```
1 z = forward_substitution(L, b_manual)
```

Finally we backward substitute:

```
x = ▶ [1.0, 2.0, 0.0]
```

```
1 x = backward_substitution(U, z)
```

and check the result:

```
▶ [0.0, 0.0, 0.0]
```

```
1 A * x - b
```

Hooray ! This succeeded!

Unfortunately **this simple algorithm does not always work**.

Consider the matrix:

```
D = 3x3 Matrix{Int64}:  
  1  2  3  
  2  4  5  
  7  8  9
```

If we apply our `factorise_lu` to this matrix we obtain:

```
3x3 LowerTriangular{Float64, Matrix{Float64}}:  
1.0      .      .  
2.0     1.0     .  
7.0    -Inf     1.0
```

```
1 let  
2   L, U = factorise_lu(D)  
3   L  
4 end
```

This `-Inf` is suspicious and points to a problem in the algorithm, which we will investigate in the next section.

LU factorisation with pivoting

We stay with the problem we identified at the end of the previous section. That is factorising the matrix

$$\mathbf{D} = \begin{pmatrix} 1 & 2 & 3 \\ 2 & 4 & 5 \\ 7 & 8 & 9 \end{pmatrix}$$

Applying **Algorithm 4** the first step ($k = 1$) will zero out the first column in the second and third row by subtracting the 2 times (7 times) the first row from the second (third) row. This results in the matrices:

$$\mathbf{A}^{(2)} = \begin{pmatrix} 1 & 2 & 3 \\ 0 & 0 & -1 \\ 0 & -6 & -12 \end{pmatrix}, \quad \mathbf{L}^{(1)} = \begin{pmatrix} 1 & & \\ 2 & & \\ 7 & & \end{pmatrix}.$$

As a result the element $A^k[k, k]$ (for $k = 2$) is zero as marked `\textcolor{red}{0}` in red.

Continuing **Algorithm 4** for $k = 2$ we would thus divide by zero in the statement $L_{ik} = \frac{A_{ik}^{(k)}}{A_{kk}^{(k)}}$ — which in the introduction of a `-Inf` in the matrix `L` that is returned by the algorithm.

We run the algorithm step by step by advancing the slider:

• nstep_lu_D =  0

► ($A^k = 3 \times 3$ Matrix{Float64} :, L = 3 × 3 Matrix{Float64} :)

1.0 2.0 3.0	0.0 0.0 0.0
-------------	-------------

We observe that from step $k = 1$ and onwards our computations are numerical garbage.

Due their central role in the Gaussian elimination algorithm the elements $(A^{(k)})_{kk}$ — respectively $A^k[k, k]$ in the implementation — are usually referred to as **pivots**.

If instead of considering the matrix A we factorise the **permuted matrix**

$$\mathbf{PD} = \begin{pmatrix} 1 & 2 & 3 \\ 7 & 8 & 9 \\ 2 & 4 & 5 \end{pmatrix}$$

which is the matrix D in which the last two rows are swapped, the algorithm goes through as expected:

```
3×3 Matrix{Int64}:
```

```
1 2 3
7 8 9
2 4 5
```

```
1 P * D
```

```
1 let
2   L, U = factorise_lu(P*D)
3   @show L
4   @show L * U
5   @show L * U - P * D   # show that L * U = P * D
6 end;
```

```
L = [1.0 0.0 0.0; 7.0 1.0 0.0; 2.0 -0.0 1.0]
L * U = [1.0 2.0 3.0; 7.0 8.0 9.0; 2.0 4.0 5.0]
L * U - P * D = [0.0 0.0 0.0; 0.0 0.0 0.0; 0.0 0.0 0.0]
```

We remark that the permutation matrix $\mathbf{P}$ is given by ...

```
P = 3x3 Matrix{Int64}:
 1  0  0
 0  0  1
 0  1  0
```

```
1 P = [1 0 0;
2      0 0 1;
3      0 1 0]
```

... and exactly achieves the task of swapping the last two rows, but leaving the rest of $\mathbf{D}$ intact as we saw above.

We notice that even though $\mathbf{D}$ cannot be permuted it is possible to obtain an LU factorisation $\mathbf{LU} = \mathbf{PD}$ if we additionally allow the freedom to cleverly permute the rows of $\mathbf{D}$.

This is in fact a general result:

Theorem 1

Every non-singular matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ admits a factorisation

$$\mathbf{PA} = \mathbf{LU}$$

where $\mathbf{L}$ is lower-triangular, $\mathbf{U}$ upper-triangular and $\mathbf{P}$ is a permutation matrix.

That is to say, that while in general **factorising a matrix $\mathbf{A}$ can fail, it always succeeds** for non-singular matrices if we **give ourselves the freedom to permute the rows of $\mathbf{A}$** .

Finding a suitable $\mathbf{P}$ can be achieved by a small modification of **Algorithm 4**. Essentially this modification boils down to *selecting* a permutation of rows of the factorised matrix on the fly, such that the pivots of the permuted matrix $(\mathbf{PA}^{(k)})_{kk}$ are always nonzero. The precise way how this is achieved is called **pivoting strategy** and the details are beyond the scope of this course. The interested reader can find some discussion in the optional subsection below.

Note, that Julia's implementation of LU factorisation (the `lu` Julia function) does indeed implement one such pivoting strategies, such that it works flawlessly on the problematic matrix $\mathbf{D}$:

```
1 facD = lu(D);
```

Notably beyond the factors $\mathbf{L}$ and $\mathbf{U}$

```
3x3 Matrix{Float64}:
 1.0  0.0  0.0
 0.285714  1.0  0.0
 0.142857  0.5  1.0
```

```
1 facD.L
```

```
3x3 Matrix{Float64}:
 7.0  8.0  9.0
 0.0  1.71429  2.42857
 0.0  0.0  0.5
```

```
1 facD.U
```

this factorisation object also contains the employed permutation matrix **P**:

```
3x3 Matrix{Float64}:
 0.0  0.0  1.0
 0.0  1.0  0.0
 1.0  0.0  0.0
```

```
1 facD.P
```

Such that as expected $\mathbf{L} * \mathbf{U} = \mathbf{P} * \mathbf{D}$:

```
3x3 Matrix{Float64}:
 0.0  0.0  0.0
 0.0  0.0  0.0
 0.0  0.0  0.0
```

```
1 facD.L * facD.U - facD.P * D
```

Solving linear systems based on LU factorisation

The result of Theorem 1 is clearly that a factorisation $\mathbf{PA} = \mathbf{LU}$ can always be achieved if the linear system $\mathbf{Ax} = \mathbf{b}$ has a unique solution, that is that **A** is non-singular. To employ pivoted LU factorisation to solve this linear system we note that we can always multiply both left and right hand sides by **P**, therefore:

$$\mathbf{Ax} = \mathbf{b} \iff \mathbf{PAx} = \mathbf{Pb} \iff \mathbf{LUx} = \mathbf{Pb}$$

which leads to the following algorithm:

Algorithm 5: Solving linear systems with pivoted LU factorisation

Given a matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$, a right-hand side $\mathbf{b} \in \mathbb{R}^n$ the linear system $\mathbf{Ax} = \mathbf{b}$ can be solved for $\mathbf{x} \in \mathbb{R}^n$ as follows:

1. Factorise $\mathbf{PA} = \mathbf{LU}$ (The `lu` Julia function).
2. Solve $\mathbf{Lz} = \mathbf{Pb}$ for $\mathbf{z}$ using *forward* substitution.
3. Solve $\mathbf{Ux} = \mathbf{z}$ for $\mathbf{x}$ using *backward* substitution.

Let us understand this algorithm by executing the steps manually for solving the problem

$$\mathbf{Dx}_D = \mathbf{b} \quad \text{with} \quad \mathbf{D} = \begin{pmatrix} 1 & 2 & 3 \\ 2 & 4 & 5 \\ 7 & 8 & 9 \end{pmatrix} \quad \text{and} \quad \mathbf{b} = \begin{pmatrix} 2 \\ 4 \\ -2 \end{pmatrix}$$

We already defined $\mathbf{D}$ and $\mathbf{b}$ before:

```
3×3 Matrix{Int64}:
```

```
1 2 3
2 4 5
7 8 9
```

```
1 D
```

```
► [2, 4, -2]
```

```
1 b
```

Step 1: Perform LU factorisation of $\mathbf{D}$:

```
► (P = 3×3 Matrix{Float64}:, L = 3×3 Matrix{Float64}:, U = 3×3 Matrix{Float64}:)
   0.0  0.0  1.0      1.0  0.0  0.0      7.0  8.0  9.0
```

```
1 begin
2   fac = lu(D)
3   (; fac.P, fac.L, fac.U)
4 end
```

Step 2: Compute $\mathbf{z}$ by forward substitution of $\mathbf{Pb}$ wrt. $\mathbf{L}$

```
zD = ► [-2.0, 4.57143, 0.0]
```

```
1 zD = forward_substitution(fac.L, fac.P * b)
```

Step 3: Backward-substitute $\mathbf{z}$ wrt. $\mathbf{U}$:

```
xD = ► [-3.33333, 2.66667, 0.0]
```

```
1 xD = backward_substitution(fac.U, zD)
```

Verify result: This solves the problem $\mathbf{Dx}_D = \mathbf{b}$ as desired:

► [0.0, 0.0, 0.0]

1 $\mathbf{D} * \mathbf{xD} - \mathbf{b}$

Conclusion: What backslash \ does to solve

To conclude this discussion we return to our starting question: What does the `\` operator do in Julia.

- We already saw that Julia essentially solves $\mathbf{A} \setminus \mathbf{b}$ by replacing it by `lu(A) \ b`.
- Julia's `lu` by default also performs a row-pivoted LU factorisation, providing thus a factorisation $\mathbf{PA} = \mathbf{LU}$.
- The returned object (an `LinearAlgebra.LU` object) internally stores both the $\mathbf{L}$, the $\mathbf{U}$ and the permutation matrix $\mathbf{P}$.
- When we use `\` with an `LU` object it, Julia immediately exploits the factorised form and performs first (permuted) forward substitution (step 2.) then backward substitution (step 3.).
- Therefore overall $\mathbf{A} \setminus \mathbf{b}$ exactly performs Algorithm 5 above.

Optional: More details on pivoting

In this section we provide some details of one common form of pivoting from LU factorisation, namely **row pivoting**.

In this approach we allow ourselves some flexibility in **Algorithm 4** by allowing ourselves to change the order in the last few rows and thus *choose* the pivot amongst the entries $\mathbf{A}_{ik}^{(k)}$ with $i \geq k$. The resulting change of row order will then define the permutation matrix $\mathbf{P}$ we employed in our above discussion.

As a guiding example we again consider the problem of factorising

$$\mathbf{D} = \begin{pmatrix} 1 & 2 & 3 \\ 2 & 4 & 5 \\ 7 & 8 & 9 \end{pmatrix}$$

where we saw previously non-pivoted LU factorisation to fail. Without any pivoting / row swapping after the first LU factorisation step (i.e after $k = 1$ is completed) the situation was

$$\mathbf{A}^{(2)} = \begin{pmatrix} 1 & 2 & 3 \\ 0 & \mathbf{0} & -1 \\ 0 & -6 & -12 \end{pmatrix}, \quad \mathbf{L}^{(1)} = \begin{pmatrix} 1 & & \\ 2 & & \\ 7 & & \end{pmatrix}.$$

Looking at $\mathbf{A}^{(2)}$ it seems very reasonable to just swap the second and the third column in $\mathbf{A}^{(2)}$ and thus move the $\mathbf{-6}$ to become the new pivot. For consistency we not only need to swap $\mathbf{A}^{(2)}$, but also $\mathbf{L}^{(2)}$. This gives us

$$\tilde{\mathbf{A}}^{(2)} = \begin{pmatrix} 1 & 2 & 3 \\ 0 & \mathbf{-6} & -12 \\ 0 & 0 & -1 \end{pmatrix}, \quad \tilde{\mathbf{L}}^{(1)} = \begin{pmatrix} 1 & & \\ 7 & & \\ 2 & & \end{pmatrix}.$$

If we continue the algorithm now, the $\mathbf{-6}$ sits in the position of the pivot $(\mathbf{A}^{(k)})_{kk}$ and the division by zero is avoided.

In fact in this particular case there is nothing left to do and the matrix $\tilde{\mathbf{A}}^{(2)}$ already is in upper triangular form. We obtain from our algorithm

$$\mathbf{U} = \mathbf{A}^{(3)} = \tilde{\mathbf{A}}^{(2)} = \begin{pmatrix} 1 & 2 & 3 \\ 0 & \mathbf{-6} & -12 \\ 0 & 0 & -1 \end{pmatrix} \quad \mathbf{L} = \mathbf{L}^{(3)} = \begin{pmatrix} 1 & 0 & 0 \\ 7 & 1 & 0 \\ 2 & 0 & 1 \end{pmatrix}.$$

Since the additional row permutation we performed not taken into account in the $\mathbf{L}$ and $\mathbf{U}$ factors, multiplying them out will not yield $\mathbf{A}$, but a matrix where the second and third rows of $\mathbf{A}$ are swapped:

$$\mathbf{LU} = \begin{pmatrix} 1 & 0 & 0 \\ 7 & 1 & 0 \\ 2 & 0 & 1 \end{pmatrix} \begin{pmatrix} 1 & 2 & 3 \\ 0 & \mathbf{-6} & -12 \\ 0 & 0 & -1 \end{pmatrix} = \begin{pmatrix} 1 & 2 & 3 \\ 7 & 8 & 9 \\ 2 & 4 & 5 \end{pmatrix}$$

To resolve this, we introduce the **permutation matrix**

$$\mathbf{P} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 1 & 0 \end{pmatrix},$$

which is obtained by performing exactly the same row permutations we did during the algorithm to the identity matrix — i.e. we swap the second and third row. With this matrix we can easily check that

$$\mathbf{LU} = \mathbf{PA}.$$

If we thus extend our `factorise_lu` function to additionally perform such pivoting permutations, the Gaussian elimination algorithm would terminate successfully.

But pivoting brings additional opportunities. As it turns out numerical stability of LU factorisation improves if one permutes not only if a pivot $(A^{(k)})_{kk}$ is zero, but in general if one *always* swaps the order of the rows such that the pivot is taken as large as possible. In other words in the k -th step of LU factorisation we always exchange row k with the row l where l satisfies

$$\left| (A^{(k)})_{lk} \right| \leq \left| (A^{(k)})_{ik} \right| \quad \text{for all } i = k, \dots, n.$$

The appropriate swaps are tracked and returned as well. Note that instead of returning a permutation matrix $\mathbf{P}$ (which requires to store n^2 elements) it is usually more convenient to store a permutation vector $\mathbf{p}$, which has only n elements. This vector tracks the indices of the rows of $\mathbf{A}$ in the order they are used as pivots. In other words if

```
idmx = 4x4 Matrix{Float64}:
 1.0  0.0  0.0  0.0
 0.0  1.0  0.0  0.0
 0.0  0.0  1.0  0.0
 0.0  0.0  0.0  1.0
```

```
1 idmx = diagm(ones(4))
```

is the identity matrix and

```
perm = ▶ [1, 3, 2, 4]
```

```
1 perm = [1, 3, 2, 4]
```

the permutation vector, then

```
4x4 Matrix{Float64}:
 1.0  0.0  0.0  0.0
 0.0  0.0  1.0  0.0
 0.0  1.0  0.0  0.0
 0.0  0.0  0.0  1.0
```

```
1 idmx[perm, :]
```

is the permutation matrix.

The code below presents an implementation of row-pivoted LU factorisation.

factorise_lu_pivot (generic function with 1 method)

```
1 function factorise_lu_pivot(A)
2     n = size(A, 1)
3     L = zeros(n, n)
4     U = zeros(n, n)
5     p = fill(0, n)
6     Ak = float(copy(A)) # Make a copy of A and ensure that all entries
7                          # are converted to floating-point numbers
8
9     for k in 1:n-1 # Algorithm steps
10         p[k] = argmax(abs.(Ak[:, k])) # Find row with maximal pivot
11
12         U[k, :] = Ak[p[k], :] # Copy pivot row to U, use U now instead of Ak,
13                               # which is again updated in-place
14         for i in 1:n # Row loop: Note the full range as any row may
15                     # be non-zero
16             L[i, k] = Ak[i, k] / U[k, k]
17             for j = 1:n # Column loop: Again full range
18                 Ak[i, j] = Ak[i, j] - L[i, k] * U[k, j]
19             end
20         end
21     end
22     p[n] = argmax(abs.(Ak[:, n]))
23     U[n, n] = Ak[p[n], n]
24     L[:, n] = Ak[:, n] / U[n, n]
25
26     # To simplify assembling L we so far kept the rows in the same order
27     # as in A. To make the matrix upper triangular we also apply the column
28     # permutation p before returning the results.
29     LowerTriangular(L[p, :]), UpperTriangular(U), p
30 end
```

We check our implementation of pivoted LU factorisation against Julia's RowMaximum pivoting strategy, which implements the same algorithm.

```
reference = LinearAlgebra.LU{Float64, Matrix{Float64}, Vector{Int64}}
L factor:
3×3 Matrix{Float64}:
 1.0      0.0  0.0
 0.285714  1.0  0.0
 0.142857  0.5  1.0
U factor:
3×3 Matrix{Float64}:
 7.0  8.0  9.0
 0.0  1.71429  2.42857
 0.0  0.0  0.5
```

```
1 reference = lu(D, RowMaximum())
```

The resulting L, U and pivot vector values are:

```
3x3 Matrix{Float64}:
 1.0  0.0  0.0
 0.285714  1.0  0.0
 0.142857  0.5  1.0
```

```
1 fac.L
```

```
3x3 Matrix{Float64}:
 7.0  8.0  9.0
 0.0  1.71429  2.42857
 0.0  0.0  0.5
```

```
1 fac.U
```

```
► [3, 2, 1]
```

```
1 fac.p
```

In contrast we obtain:

```
3x3 LowerTriangular{Float64, Matrix{Float64}}:
 1.0  .  .
 0.285714  1.0  .
 0.142857  0.5  1.0
```

```
1 factorise_lu_pivot(D)[1] # L
```

```
3x3 UpperTriangular{Float64, Matrix{Float64}}:
 7.0  8.0  9.0
 .  1.71429  2.42857
 .  .  0.5
```

```
1 factorise_lu_pivot(D)[2] # U
```

```
► [3, 2, 1]
```

```
1 factorise_lu_pivot(D)[3] # p
```

Memory usage and fill-in

Sparse matrices

The matrices and linear systems one encounters in physics and engineering applications not infrequently reach **huge sizes**. For example in the numerical solution of partial differential equations using finite difference or finite element methods the vector of unknowns frequently has a **dimension on the order of $n \simeq 10^6$** .

Taking $n = 10^6$ thus as an example this implies that matrices, i.e. entities from $\mathbb{R}^{n \times n}$, hold around 10^{12} **elements**. In double precision, i.e. **64-bit** numbers, each entry requires **8 bytes**. As a result **storing all elements** of such a matrix explicitly in memory requires

```
8000000000000
```

```
1 10^12 * 8
```

bytes or

```
7.275957614183426
```

```
1 10^12 * 8 / 1024^4 # TiB
```

tibibytes of storage. **This is a challenge** even for modern compute clusters, where typically a node has around **512** GiB. Laptops nowadays feature around **32** GiB, so would be completely out of the question to solve such problems.

However, **in many applications** the arising **matrices are sparse**, meaning that they contain a large number of zero elements, which do not need to be stored. Let us discuss a few examples.

Definition: Full matrix

A matrix $A \in \mathbb{R}^{n \times n}$ is called **full** if the number of non-zero elements is at the order of n^2 . For such matrices almost all elements are non-zero and need to be stored.

Full matrices are the "standard case" and for them memory constraints usually set the limit of the size of problems, which can be tackled. For example an **32** GiB memory laptop can store around $4 \cdot 10^9$ double-precision numbers, which means that linear problems with more than around $n \simeq 60000$ unknowns cannot be treated.


Definition: Sparse matrix

A matrix $A \in \mathbb{R}^{n \times n}$ is called **sparse** if the number of non-zero elements is at the order of n .

If we know which elements are zero, we can thus save memory by only storing those elements, which are non-zero. For this purpose the [SparseArrays](#) [Julia package](#) implements many primitives for working with sparse arrays.

This includes, generating random sparse arrays:

`Asp = 100×100 SparseMatrixCSC{Float64, Int64} with 1541 stored entries:`



```
1 Asp = sprand(100, 100, 0.15)
```

or simply sparsifying a dense array using the `sparse` function, which converts a full matrix to a sparse matrix by dropping the explicit storage of all zero entries.

`M = 4×4 Matrix{Int64}:`

```
 1  0  0  5
-2  0  1  0
 0  0  6  0
 0  1  0 -1
```

```
1 M = [ 1 0 0 5;
2      -2 0 1 0;
3        0 0 6 0;
4        0 1 0 -1]
```

`4×4 SparseMatrixCSC{Int64, Int64} with 7 stored entries:`

```
 1  .  .  5
-2  .  1  .
 .  .  6  .
 .  1  . -1
```

```
1 sparse(M)
```

Using the `SparseArray` data structure from `SparseArrays` consistently allows to fully exploit sparsity when solving a problem. As a result the **storage costs scale only as $O(n)$** . With our laptop of 32 GiB memory we can thus tackle problems with around $n \simeq 4 \cdot 10^9$ unknowns — much better than the **60000** we found when using full matrices.

Finally we introduce a special kind of sparse matrix:


Definition: Band matrix

A matrix $A \in \mathbb{R}^{n \times n}$ is called a **band matrix** with bandwidth d if $A_{ij} = 0$ when $|j - i| > d$. Every line of the matrix contains at most $2d + 1$ non-zero elements and the number of non-

zeros thus scales as $O(nd)$.

An example for a banded matrix with bandwidth **5** is:

`band = 105×105 SparseMatrixCSC{Float64, Int64} with 817 stored entries:`



Memory: LU factorisation of full matrices

Since the amount of available memory can put hard constraints on the size of linear systems which can be solved, we now want to investigate the memory requirement of LU factorisation $\mathbf{A} = \mathbf{L}\mathbf{U}$.

If $\mathbf{A}$ is a **full matrix**, then we know that $\mathbf{L}$ and $\mathbf{U}$ are triangular, thus $\mathbf{L}$ and $\mathbf{U}$ contain less non-zero elements or **more zero elements** than $\mathbf{A}$ itself and thus **require together as much memory** to be stored in memory as $\mathbf{A}$ itself.


Memory: LU factorisation of sparse matrices

Let's **contrast** this **with the sparse matrices** we have considered above. First the random sparse matrix:

`100×100 SparseMatrixCSC{Float64, Int64} with 1541 stored entries:`



100×100 SparseMatrixCSC{Float64, Int64} with 3321 stored entries:



```
1 lu(Asp).L # U looks similar
```

Storing both L and U thus requires a total number of


6475

```
1 nnz(lu(Asp).L) + nnz(lu(Asp).U)
```

non-zero elements, which is about **4.2 times as much** as the original matrix !


Now the banded matrix:

105×105 SparseMatrixCSC{Float64, Int64} with 817 stored entries:



```
1 band
```

105×105 SparseMatrixCSC{Float64, Int64} with 611 stored entries:



```
1 lu(band).L # U looks similar
```

At least the banded structure is kept, but still we require

1222

```
1 nnz(lu(band).L) + nnz(lu(band).U)
```

nonzeros, which is about 1.5 times than the original in this case.

In both cases storing the **L** and the **U** factors require more non-zero elements than the original matrix. This phenomenon is usually referred to as **fill in**.

In particular for the sparse matrix `Asp` LU factorisation did not preserve the structure at all. In contrast it almost resulted in a full matrix in the lower right corner of the **L**. For **U** this looks exactly the same. In fact one can show that in general even for sparse matrices the **memory usage of LU factorisation** is $O(n^2)$. Therefore while one may be able to store a sparse matrix **A** with a huge size in memory (due to the $O(n)$ memory cost), one may actually run out of memory while computing the factorisation.

Let us add that for banded matrices the situation is slightly better as one can show that in this case the fill-in takes place at most inside the band. As a result the memory requirement stays at $O(nd)$.

Overview of LU factorisation memory cost

We summarise the cost of LU factorisation in a table:

type of $n \times n$ matrix	memory usage	comment
full matrix	$O(n^2)$	
general sparse matrix	$O(n^2)$	fill-in

Computational cost of LU factorisation

In this section we want to investigate how long it will take a computer to perform LU factorisation on a large matrix.

Modern laptop computers nowadays have a clock frequency of a few Gigahertz (GHz). This means they are **able to perform about 10^9 operations per second**, where for simplicity we assume that "one operation" is an elementary addition, multiplication, division and so on. If we can determine how many such operations are needed to perform LU factorisation we can estimate the computational cost.

Before we look at the LU algorithm (Algorithm 4) we first understand a few simpler cases from linear algebra:

Scalar product

Given two vectors $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$ consider computing the scalar product

$$\mathbf{x} \cdot \mathbf{y} = \mathbf{x}^T \mathbf{y} = \sum_{i=1}^n x_i y_i$$

which in code is achieved as

scalar_product (generic function with 1 method)

```
1 function scalar_product(x, y)
2     result = 0.0
3     for i in 1:length(x)
4         result = result + x[i] * y[i]
5     end
6     result
7 end
```

In the loop for each iteration we require **1** multiplication and **1** addition. In total the scalar_product function therefore requires n multiplications and n additions. The number of elementary operations is thus $2n$. We say the **computational cost is $O(n)$** (i.e. on the order of n or linear in n).

If we take the dimensionality $n = 1000$ the number of operations is $O(1000)$. On a 1 GHz computer (with 10^9 operations per second) this therefore takes about 10^{-6} seconds ... hardly

noticable.

Matrix-vector product

Given a matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ and a vector $\mathbf{x} \in \mathbb{R}^n$ the matrix-vector product $\mathbf{y} = \mathbf{A}\mathbf{x}$ is computed as

$$y_i = \sum_{j=1}^n A_{ij}x_j \quad \text{for } i = 1, \dots, n$$

or in code

matrix_vector_product (generic function with 1 method)

```
1 function matrix_vector_product(A, x)
2     y = zeros(size(A, 1)) # Create a vector of zeros (#)
3
4     for i in 1:size(A, 1)
5         for j in 1:size(A, 2)
6             y[i] = y[i] + A[i, j] * x[j] # (*)
7         end
8     end
9     result
10 end
```

In the innermost loop we observe again that we require **1** addition and **1** multiplication per iteration. This instruction is performed once for each combination of i and j , so we look at the limits of each of the nested loops. The loop over i runs over n values (the number of rows of $\mathbf{A}$) and the loop over j over n values as well (the number of columns of $\mathbf{A}$). In total the inner most instruction $(*)$ is thus run n^2 times, each times costing **2** operations. The total **cost is thus** $O(n^2)$.

For our case with $n = 1000$ and a 1 GHz computer we thus now need $(10^3)^2/10^9 s = 10^{-3} s = 1\text{ms}$, which is again a rather short time.

► **Expert information:** What about the allocation in ``(#)`` ?

Overview of computational cost

For vectors $\mathbf{v}, \mathbf{w} \in \mathbb{R}^n$ and matrices $\mathbf{A} \in \mathbb{R}^{n \times n}$ the computational cost is

operation	cost
dot product $\mathbf{v}^T \mathbf{w}$	$O(n^2)$

Finally a few rough guidelines:

General guideline to estimate computational cost

1. Determine the instructions at the innermost (most deeply nested) loop level
2. Find the most expensive of these instructions and determine its cost in terms of *number of operations*
3. Determine the ranges of all loops in terms of the dimensionality of your problem. Typically for each loop level this is n
4. Multiply the result of 2. and all index ranges of 3 to get the total scaling. Typically for a single loop nesting the cost is $O(n)$ for a doubly nested loop $O(n^2)$ and so on.

LU factorisation

We revisit Algorithm 4:

Algorithm 4: LU factorisation

Input: $\mathbf{A} \in \mathbb{R}^{n \times n}$

Output: $\mathbf{U} \in \mathbb{R}^{n \times n}$, $\mathbf{L} \in \mathbb{R}^{n \times n}$

- $\mathbf{A}^{(1)} = \mathbf{A}$
- for $k = 1, \dots, n - 1$ (algorithm steps)
 - $L_{kk} = 1$
 - for $i = k + 1, \dots, n$ (Loop over rows)
 - $L_{ik} = \frac{A_{ik}^{(k)}}{A_{kk}^{(k)}}$
 - for $j = k, \dots, n$ (Loop over columns)
 - $A_{ij}^{(k+1)} = A_{ij}^{(k)} - L_{ik}A_{kj}^{(k)}$
- $L_{nn} = 1$ (the loop above only runs until $n - 1$)
- $\mathbf{U} = \mathbf{A}^{(n)}$

We notice that the instruction at the innermost loop level is

$$A_{ij}^{(k+1)} = A_{ij}^{(k)} - L_{ik}A_{kj}^{(k)},$$

which costs again 1 addition and 1 multiplication and which is executed once for each tuple (i, j, k) . In the worst case k takes $n - 1$ elements, i also $n - 1$ elements (assume $k = 1$) and j again in its worst case takes $n - 1$ elements. In total the cost is thus $(n - 1)^3 = O(n^3)$.

For our case with $n = 1000$ and a 1 GHz computer the computation thus needs approximately $(10^3)^3/10^9 s = 1s$, which starts to be a noticable amount of time. If we even consider $n = 10^5$ all of a sudden it takes $10^3 s$, which is about 15 minutes. For **large matrices the cost of computing LU factorisation can thus become important !**

For **banded matrices** the computational scaling is a little improved. E.g. consider LU factorisation on a banded matrix A with band width d . Then both the loop over i (row loop) as well as the loop over j (column loop) in the LU factorisation algorithm can be truncated and at most run over d elements. As a result the computational cost at worst becomes $O(nd^2)$ which for a small band width (i.e. $d < n$) is substantially less than $O(n^3)$.

Overview of LU factorisation computational cost and memory cost

We summarise the cost of LU factorisation in a table:

type of $n \times n$ matrix	computational cost	memory usage
full matrix	$O(n^3)$	$O(n^2)$
general sparse matrix		
banded matrix, band width d	$O(nd^2)$	$O(nd)$

Numerical stability

To close off our discussion of Gaussian elimination we consider the question of its numerical stability, i.e. how is the **result of LU factorisation effected by small round-off errors** introduced when performing a calculation on a computer with its finite precision to represent numbers.

In this chapter our goal has been to solve a linear system of the form

$$\mathbf{Ax} = \mathbf{b}, \tag{5}$$

where we are given a matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ as well as a right-hand side $\mathbf{b} \in \mathbb{R}^n$ and we wish to solve for $\mathbf{x}$. We will later refer to this system as the **exact system**.

Since the computer in general is unable to represent the matrix and right-hand side exactly in finite-precision floating point arithmetic, even **just providing this problem to a computer** will

usually be associated with **making a small error**: in practice the computer only ever sees the **perturbed system**

$$\tilde{\mathbf{A}}\tilde{\mathbf{x}} = \tilde{\mathbf{b}}, \quad (6)$$

where $\tilde{\mathbf{A}}$ is an approximation to $\mathbf{A}$, $\tilde{\mathbf{b}}$ is an approximation to $\mathbf{b}$. It will thus yield the solution $\tilde{\mathbf{x}}$ by solving the perturbed system instead of providing $\mathbf{x}$.

To understand the numerical stability when solving linear systems our goal is thus to understand how far the solution $\tilde{\mathbf{x}}$ — obtained on a computer — differs from $\mathbf{x}$ — the solution to (5), the true linear system we want to solve.

To provide a motivation **why this type of analysis matters in practice**, we first consider the following example:

Example: Rounding error in a 2x2 system

Consider the *exact* linear system

$$\mathbf{A}\mathbf{x} = \mathbf{b} \quad \Leftrightarrow \quad \begin{pmatrix} 1 & 10^{-16} \\ 1 & 0 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \begin{pmatrix} 1 \\ 1 \end{pmatrix}$$

with solution $(x_1, x_2) = (1, 0)$ and the *perturbed* linear system

$$\mathbf{A}\tilde{\mathbf{x}} = \tilde{\mathbf{b}} \quad \Leftrightarrow \quad \begin{pmatrix} 1 & 10^{-16} \\ 1 & 0 \end{pmatrix} \begin{pmatrix} \tilde{x}_1 \\ \tilde{x}_2 \end{pmatrix} = \begin{pmatrix} 1 + 10^{-16} \\ 1 \end{pmatrix},$$

where only the first component of the right-hand side has been perturbed by 10^{-16} . This second system has solution $(\tilde{x}_1, \tilde{x}_2) = (1, 1)$. The second component of this solution is thus *completely wrong*!

The small perturbation of 10^{-16} , which may readily occur just by inputting the data to a computer, can already have a very noticable effect on the solution and in fact render the solution we obtain from our computation (i.e. $\tilde{\mathbf{x}}$) rather far from the answer we are actually after (i.e. $\mathbf{x}$).

In standard floating-point arithmetic the **relative error** in representing any number is around 10^{-16} . Mathematically we can find the following relationships between the elements of $\tilde{\mathbf{b}}$ and $\mathbf{b}$ as well as $\tilde{\mathbf{A}}$ and $\mathbf{A}$, respectively:

$$\begin{aligned}\tilde{b}_i &= b_i (1 + \varepsilon_i) & \text{where } \varepsilon_i \text{ is on the order of } 10^{-16} \\ \tilde{A}_{ij} &= A_{ij} (1 + \eta_{ij}) & \text{where } \eta_{ij} \text{ is on the order of } 10^{-16}\end{aligned}$$

However, to simplify the treatment in this course we will not attempt to discuss the effect of such relative errors on the numerical stability in full generality. Much rather we will employ the **error model**:

Error model in this notebook

Between the exact system and the perturbed system we will assume the foollowing relationship

$$\begin{aligned}\tilde{b}_i &= b_i (1 + \varepsilon) & \text{where } \varepsilon \approx 10^{-16} \\ \tilde{A}_{ij} &= A_{ij}\end{aligned}$$

In other words we assume that the system matrix **A** is exactly represented on the computer and that round-off errors only affect the right-hand side **$\tilde{\mathbf{b}}$** . Moreover we assume that the relative error for all elements of **b** is identical to ε .

For standard double-precision floating-point numbers we have $\epsilon \approx 10^{-16}$.

For an alternative and more detailed discussion see also [chapter 2.8](#) of Driscoll, Brown: *Fundamentals of Numerical Computation*.

Matrix condition numbers and stability result

To analyse the error $\mathbf{x} - \tilde{\mathbf{x}}$ mathematically, we first have to introduce some notation.

Recall that for a vector $\mathbf{v} \in \mathbb{R}^n$ its Euclidean norm is $\|\mathbf{v}\| = \sqrt{\sum_{i=1}^n v_i^2}$.

We define the following:

Definition: Relative error of linear systems

The **relative error** in $\tilde{\mathbf{x}}$, the solution to the perturbed system (6), is the quantity

$$\frac{\|\mathbf{x} - \tilde{\mathbf{x}}\|}{\|\mathbf{x}\|},$$

where $\mathbf{x}$ is the exact solution, i.e. the solution to the exact linear system (5) one actually wanted to solve.

We further need a generalisation of norms to matrices:

Definition: Matrix norm

Given $\mathbf{M} \in \mathbb{R}^{m \times n}$ a real matrix (not necessarily square), we define the matrix norm of $\mathbf{M}$ as

$$\|\mathbf{M}\| = \max_{\substack{\mathbf{v} \in \mathbb{R}^n \\ \mathbf{v} \neq \mathbf{0}}} \frac{\|\mathbf{M}\mathbf{v}\|}{\|\mathbf{v}\|}.$$

The previous definition implies in particular that

$$\|\mathbf{M}\mathbf{v}\| \leq \|\mathbf{M}\| \|\mathbf{v}\| \quad \forall \mathbf{x} \in \mathbb{R}^n. \quad (7)$$

With this definitions in place we begin our analysis for (5) and (6) where $\mathbf{A} \in \mathbb{R}^{n \times n}$. From the definition of the exact and perturbed linear systems

$$\mathbf{A}\mathbf{x} = \mathbf{b} \quad \text{and} \quad \mathbf{A}\tilde{\mathbf{x}} = \tilde{\mathbf{b}}$$

we obtain by subtraction and since $\mathbf{A}$ is invertible

$$\mathbf{A}(\mathbf{x} - \tilde{\mathbf{x}}) = \mathbf{b} - \tilde{\mathbf{b}} \quad \Rightarrow \quad \mathbf{x} - \tilde{\mathbf{x}} = \mathbf{A}^{-1}(\mathbf{b} - \tilde{\mathbf{b}})$$

Using (7) we therefore have

$$\|\mathbf{x} - \tilde{\mathbf{x}}\| \leq \|\mathbf{A}^{-1}\| \|\mathbf{b} - \tilde{\mathbf{b}}\|.$$

Furthermore we have from applying (7) to $\mathbf{A}\mathbf{x} = \mathbf{b}$:

$$\|\mathbf{b}\| \leq \|\mathbf{A}\| \|\mathbf{x}\| \quad \Rightarrow \quad \frac{1}{\|\mathbf{x}\|} \leq \frac{\|\mathbf{A}\|}{\|\mathbf{b}\|}$$

such that with the above result we can bound the relative error as

$$\frac{\|\mathbf{x} - \tilde{\mathbf{x}}\|}{\|\mathbf{x}\|} \leq \|\mathbf{A}^{-1}\| \|\mathbf{b} - \tilde{\mathbf{b}}\| \cdot \frac{\|\mathbf{A}\|}{\|\mathbf{b}\|} = \underbrace{\|\mathbf{A}^{-1}\| \|\mathbf{A}\|}_{=\kappa(\mathbf{A})} \frac{\|\mathbf{b} - \tilde{\mathbf{b}}\|}{\|\mathbf{b}\|}. \quad (8)$$

From a stability analysis point of view the quantity $\kappa(\mathbf{A}) = \|\mathbf{A}^{-1}\| \|\mathbf{A}\|$ thus relates the relative error in the right-hand side (input quantity) to the relative error in the solution $\mathbf{x}$ (output quantity). It is thus the condition number for a linear system:

Definition: Condition number of a linear system

Given a linear system $\mathbf{Ax} = \mathbf{b}$ with an invertible system matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ its **condition number** is defined as

$$\kappa(\mathbf{A}) = \|\mathbf{A}^{-1}\| \|\mathbf{A}\|.$$

Due to the importance of solving linear systems in linear algebra one usually calls $\kappa(\mathbf{A})$ also the **condition number of the matrix $\mathbf{A}$** .

But notably (8) can be interpreted even more generally as the way in which the solution of linear systems are changed as we change the right-hand side as is summarised below:

Theorem 2: Change of solution when changing the right-hand side

Given a linear system $\mathbf{Ax} = \mathbf{b}$, which we solve for $\mathbf{x}$ and a related linear system $\mathbf{A}\tilde{\mathbf{x}} = \tilde{\mathbf{b}}$ which we solve for $\tilde{\mathbf{x}}$, then the two solutions are related as

$$\frac{\|\mathbf{x} - \tilde{\mathbf{x}}\|}{\|\mathbf{x}\|} \leq \kappa(\mathbf{A}) \frac{\|\mathbf{b} - \tilde{\mathbf{b}}\|}{\|\mathbf{b}\|}.$$

In this notebook we employed an error model where storing $\mathbf{b}$ on the computer (which leads to $\tilde{\mathbf{b}}$) introduces a small relative error:

$$\tilde{b}_i = b_i (1 + \varepsilon).$$

This enables to simplify the expression of the absolute error of right-hand side:

$$\|\mathbf{b} - \tilde{\mathbf{b}}\| = \sqrt{\sum_{i=1}^n b_i^2 \varepsilon^2} = \varepsilon \sqrt{\sum_{i=1}^n b_i^2} = \varepsilon \|\mathbf{b}\|,$$

which leads to the following result:

Theorem 3: Stability of solving linear systems

Given a linear system $\mathbf{Ax} = \mathbf{b}$ and $\tilde{\mathbf{x}}$ the solution to a perturbed linear system $\mathbf{A}\tilde{\mathbf{x}} = \tilde{\mathbf{b}}$, where $\tilde{b}_i = b_i(1 + \epsilon)$ for $i = 1, \dots, n$, then

$$\frac{\|\mathbf{x} - \tilde{\mathbf{x}}\|}{\|\mathbf{x}\|} \leq \kappa(\mathbf{A}) \epsilon$$

This inequality shows that the condition number of the matrix $\mathbf{A}$ plays the role of *amplifying* the round-off error introduced by the floating-point representation of the right-hand side.

Since for standard double-precision floating-point numbers we have $\epsilon \approx 10^{-16}$, this means that condition numbers of 10^{16} lead to a relative error of 1, i.e. 100% error — all precision is lost. In general condition numbers above 10^8 start to become problematic as in this case the relative error in the solution is at most 10^{-8} , i.e. roughly speaking no more than 8 digits.

Numerically computing the condition number. In Julia code the condition number of a matrix is computed using the `cond` function. It effectively uses the above formulas for its computation.

For example for our running example we obtain:

```
2.0000000000000004e16
```

```
1 let
2     A = [1 1e-16;
3          1 0 ]
4     κ = cond(A)
5 end
```

Notice, that this is $2 \cdot 10^{16}$, which is a huge number !

Rounding error in a 2x2 system (continued)

Using Theorem 3 we can finally understand the behaviour we observe in our example. Previously we found that a small perturbation $\mathbf{1} \rightarrow \mathbf{1} + 10^{-16}$ in one of the elements of the right hand side, introduces a change in the solution from $(x_1, x_2) = (1, 0)$ to $(\tilde{x}_1, \tilde{x}_2) = (1, 1)$. This is a relative error of

$$\frac{\|\mathbf{x} - \tilde{\mathbf{x}}\|}{\|\mathbf{x}\|} = \frac{1}{1} = 1.$$

Note that in this example $\epsilon \approx 10^{-16}$ just like for the case of a rounding error in the right-hand side.

As we computed above

$$\kappa(\mathbf{A}) \approx 2 \cdot 10^{16},$$

such that this large relative error is explained by the large condition number of this matrix.

Computing condition numbers

Using the `cond` function of Julia enables us to easily compute condition numbers in practice.

However, the expression $\kappa(\mathbf{A}) = \|\mathbf{A}^{-1}\| \|\mathbf{A}\|$ is not extremely handy to provide a good intuition for what the condition number actually means or how it varies as the matrix $\mathbf{A}$ is changed.

In this section we thus discuss a few standard techniques how to compute condition numbers of matrices. These techniques are in fact exactly the algorithms that `cond` uses under the hood to do its computation.

First we need some notation:

Definition: Minimal and maximal eigenvalues

Given a diagonalisable matrix $\mathbf{M} \in \mathbb{R}^{n \times n}$ we denote by $\lambda_{\max}(\mathbf{M})$ the *largest* eigenvalue of $\mathbf{M}$ and by $\lambda_{\min}(\mathbf{M})$ the *smallest* eigenvalue of $\mathbf{M}$, i.e.

$$\lambda_{\max}(\mathbf{M}) = \max_{i=1, \dots, n} \lambda_i(\mathbf{M}) \quad \text{and} \quad \lambda_{\min}(\mathbf{M}) = \min_{i=1, \dots, n} \lambda_i(\mathbf{M})$$

where $\lambda_i(\mathbf{M})$ for $i = 1, \dots, n$ are the eigenvalues of $\mathbf{M}$.

It turns out that the minimal and maximal eigenvalues provide a convenient way to compute matrix norms:

Lemma 4: Computing matrix norms

For any matrix $\mathbf{M} \in \mathbb{R}^{m \times n}$

$$\|\mathbf{M}\| = \sqrt{\lambda_{\max}(\mathbf{M}^T \mathbf{M})}.$$

If $\mathbf{M}$ is moreover square and invertible, then

$$\|\mathbf{M}^{-1}\| = \sqrt{\lambda_{\max}(\mathbf{M}^{-T}\mathbf{M}^{-1})} = \frac{1}{\sqrt{\lambda_{\min}(\mathbf{M}^T\mathbf{M})}}.$$

Based on this we can deduce a few useful formulas for computing condition numbers. We start with the general expression:

Corollary 5: Condition number for invertible matrices

If $\mathbf{A} \in \mathbb{R}^{n \times n}$ is invertible the condition number can be computed as

$$\kappa(\mathbf{A}) = \|\mathbf{A}^{-1}\| \|\mathbf{A}\| = \frac{\sqrt{\lambda_{\max}(\mathbf{A}^T\mathbf{A})}}{\sqrt{\lambda_{\min}(\mathbf{A}^T\mathbf{A})}}.$$

Furthermore if $\mathbf{A} \in \mathbb{R}^{n \times n}$ is invertible and **symmetric** then $\mathbf{A}^T = \mathbf{A}$. As a result we observe that

$$\lambda_i(\mathbf{A}^T\mathbf{A}) = \lambda_i(\mathbf{A}^2) = \lambda_i(\mathbf{A})^2 \quad \text{for all } i = 1, \dots, n.$$

As a result

$$\kappa(\mathbf{A}) = \frac{\sqrt{\lambda_{\max}(\mathbf{A}^T\mathbf{A})}}{\sqrt{\lambda_{\min}(\mathbf{A}^T\mathbf{A})}} = \frac{\sqrt{\max_{i=1, \dots, n} \lambda_i(\mathbf{A})^2}}{\sqrt{\min_{i=1, \dots, n} \lambda_i(\mathbf{A})^2}} = \frac{\max_{i=1, \dots, n} |\lambda_i(\mathbf{A})|}{\min_{i=1, \dots, n} |\lambda_i(\mathbf{A})|}$$

If $\mathbf{A} \in \mathbb{R}^{n \times n}$ is **both symmetric and positive definite** (i.e. all eigenvalues are strictly positive), then we can drop the moduli in the above expression and obtain:

Lemma 6: Condition number of symmetric, positive-definite matrices

If $\mathbf{A} \in \mathbb{R}^{n \times n}$ is symmetric and positive definite, then

$$\kappa(\mathbf{A}) = \frac{\lambda_{\max}(\mathbf{A})}{\lambda_{\min}(\mathbf{A})}$$

Rounding error in a 2x2 system (continued)

As an example we now compute the condition number of the matrix analytically

$$\mathbf{A} = \begin{pmatrix} 1 & 10^{-16} \\ 1 & 0 \end{pmatrix}.$$

Recall that the value computed by Julia's `cond` function was roughly $2 \cdot 10^{16}$.

This matrix is not symmetric and we therefore use the expression

$$\kappa(\mathbf{A}) = \frac{\sqrt{\lambda_{\max}(\mathbf{A}^T \mathbf{A})}}{\sqrt{\lambda_{\min}(\mathbf{A}^T \mathbf{A})}},$$

which requires us to compute the eigenvalues of $\mathbf{A}^T \mathbf{A}$. We note

$$\mathbf{A}^T \mathbf{A} = \begin{pmatrix} 1 & 1 \\ 10^{-16} & 0 \end{pmatrix} \begin{pmatrix} 1 & 10^{-16} \\ 1 & 0 \end{pmatrix} = \begin{pmatrix} 2 & 10^{-16} \\ 10^{-16} & 10^{-32} \end{pmatrix}$$

and

$$\det(\mathbf{A}^T \mathbf{A} - \lambda I) = \lambda^2 - (2 + 10^{-32})\lambda + 10^{-32},$$

such that the two eigenvalues of $\mathbf{A}^T \mathbf{A}$ are

$$\lambda_{\pm}(\mathbf{A}^T \mathbf{A}) = \frac{1}{2} \left(2 + 10^{-32} \pm \sqrt{4 + 10^{-64}} \right).$$

We conclude that

$$\lambda_{\max}(\mathbf{A}^T \mathbf{A}) \approx 2 \quad \text{and} \quad \lambda_{\min}(\mathbf{A}^T \mathbf{A}) \approx \frac{10^{-32}}{2}$$

therefore

$$\kappa(\mathbf{A}) \approx 2 \cdot 10^{16},$$

— i.e. the same huge number we obtained before.

Numerical analysis

1. Introduction
2. The Julia programming language
3. Revision and preliminaries
4. Root finding and fixed-point problems
5. Interpolation
6. Direct methods for linear systems
7. Iterative methods for linear systems

8. Eigenvalue problems

9. Numerical integration